

A series of horizontal bars of varying lengths and colors (teal, blue, green) are positioned on the left side of the slide, creating a modern, abstract background element.

Module 8: Beyond Linearity

Content based on ISLR Chapter 7.1 – 7.7

Moving Beyond Linearity

- The two types of regression functions that we have looked at (linear and logit) are surprisingly effective with many types of models and have advantages of interpretability.
- However, the underlying structure of the data is sometimes too non-linear to be effectively modeled with these equations and more complex models must be considered.

Moving Beyond Linearity

Outline

- Non-linear Regression Overview
 - Basis Functions
 - Polynomial Regression
 - Piecewise-Constant Regression Models (Step Functions)
 - Regression splines
 - Local regression
 - General Additive Models

Basis Functions

- A general approach to extending the linear regression and classification approaches involves the use of basis functions
- Instead of calculating regression coefficients on the explanatory attributes, we apply them to a function of the explanatory attributes:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \cdots + \varepsilon$$

$$Y = \beta_0 + \beta_1 b_1(X_1) + \beta_2 b_2(X_2) + \beta_3 b_3(X_3) + \cdots + \varepsilon$$

where $b_1(\cdot)$ is the basis function applied to the explanatory attribute

Polynomial Regression

- Already covered in Module 3 as a simple extension of the linear model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \beta_3 X_1^3 + \dots + \varepsilon$$

- As previously noted, this is still a linear model and we can use all of the linear regression techniques by simply creating a new transformed attributes (X_1^2, X_1^3, \dots)
- Polynomial regression is one example of the use of a basis function approach where the basis function is given as:

$$b_j(X_1) = X_1^j$$

Wage Example

Polynomial Regression

	year	age	sex	mariti	race	education	region	jobclass	health	health_ins	logwage	wage
0	2006	18	1. Male	1. Never Married	1. White	1. < HS Grad	2. Middle Atlantic	1. Industrial	1. <=Good	2. No	4.318063	75.043154
1	2004	24	1. Male	1. Never Married	1. White	4. College Grad	2. Middle Atlantic	2. Information	2. >=Very Good	2. No	4.255273	70.476020
2	2003	45	1. Male	2. Married	1. White	3. Some College	2. Middle Atlantic	1. Industrial	1. <=Good	1. Yes	4.875061	130.982177
3	2003	43	1. Male	2. Married	3. Asian	4. College Grad	2. Middle Atlantic	2. Information	2. >=Very Good	1. Yes	5.041393	154.685293
4	2005	50	1. Male	4. Divorced	1. White	2. HS Grad	2. Middle Atlantic	2. Information	1. <=Good	1. Yes	4.318063	75.043154
—	—	—	—	—	—	—	—	—	—	—	—	—
2995	2008	44	1. Male	2. Married	1. White	3. Some College	2. Middle Atlantic	1. Industrial	2. >=Very Good	1. Yes	5.041393	154.685293
2996	2007	30	1. Male	2. Married	1. White	2. HS Grad	2. Middle Atlantic	1. Industrial	2. >=Very Good	2. No	4.602060	99.689464
2997	2005	27	1. Male	2. Married	2. Black	1. < HS Grad	2. Middle Atlantic	1. Industrial	1. <=Good	2. No	4.193125	66.229408
2998	2005	27	1. Male	1. Never Married	1. White	3. Some College	2. Middle Atlantic	1. Industrial	2. >=Very Good	1. Yes	4.477121	87.981033
2999	2009	55	1. Male	5. Separated	1. White	2. HS Grad	2. Middle Atlantic	1. Industrial	1. <=Good	1. Yes	4.505150	90.481913


3000 rows × 12 columns

Wage Example

Polynomial Regression

```
1 import seaborn as sns
2 sns.set(rc = {'figure.figsize': (15,8)})
3 sns.scatterplot(data = wage, x = 'age', y = 'wage')
```

```
<AxesSubplot:xlabel='age', ylabel='wage'>
```



Wage Example

Polynomial Regression

Define sklearn polynomial regression model using Sklearn pipeline function

```
] 1 from sklearn.preprocessing import PolynomialFeatures
   2 from sklearn.linear_model import LinearRegression
   3 from sklearn.pipeline import Pipeline
   4
   5 # Pipeline chains together multiple steps
   6 poly_model = Pipeline([('poly', PolynomialFeatures(degree=4)), ('linear', LinearRegression())])
   7 poly_model.fit(wage_data[['age']], wage_data['wage'])

]: Pipeline(steps=[('poly', PolynomialFeatures(degree=4)),
                   ('linear', LinearRegression())])
```



Wage Example

Polynomial Regression

Create array of ages and use model to predict wages to create curve and plot

```
1 ages = np.arange(wage['age'].min(), wage['age'].max(), 1)
2 y_hat = poly_model.predict(ages.reshape(-1,1))
3 sns.set(rc = {'figure.figsize':(15,8)})
4 sns.scatterplot(data = wage, x = 'age', y = 'wage')
5 sns.lineplot(x=ages, y=y_hat, color = 'r')
```


<AxesSubplot:xlabel='age', ylabel='wage'>



Wage Example

Polynomial Regression


```
1 sns.regplot(x=wage['age'], y=wage['wage'], ci=None, order=4, line_kws={"color": "red"});
```



Piecewise-Constant Regression Models

"Step Functions"

- Instead of imposing a global polynomial structure on model, we cut the variable into distinct regions and fit with piecewise constants:



Step Function

- Model is easily fit by creating *Indicator Functions* $I(\cdot)$ that return 1 if the condition is true and 0 otherwise)
- To create the step function model in our example, we need three step functions:
 - $C_1(Radio) = I(Radio < 2)$
 - $C_2(Radio) = I(2 \leq Radio \leq 4)$
 - $C_3(Radio) = I(Radio > 4)$
- Then, we use standard least squares to fit a linear model with these the indicators as predictors:

$$Sales = \beta_0 + \beta_1 C_1(Radio) + \beta_2 C_2(Radio) + \beta_3 C_3(Radio)$$

Step Function

Wage Example

Step Function

```
1 df_cut, bins = pd.cut(wage['age'], 4, retbins=True, right=True)
2 bins
```

```
array([17.938, 33.5 , 49. , 64.5 , 80. ])
```

```
1 df_cut
```

```
0      (17.938, 33.5]
1      (17.938, 33.5]
2      (33.5, 49.0]
3      (33.5, 49.0]
4      (49.0, 64.5]
```

```
...
2995     (33.5, 49.0]
2996     (17.938, 33.5]
2997     (17.938, 33.5]
2998     (17.938, 33.5]
2999     (49.0, 64.5]
```

```
Name: age, Length: 3000, dtype: category
```

```
Categories (4, interval[float64]): [(17.938, 33.5] < (33.5, 49.0] < (49.0, 64.5] < (64.5, 80.0]]
```

Step Function

Wage Example

```
1 df_steps = pd.concat([wage['age'], df_cut, wage['age']], keys=['age', 'age_cuts', 'wage'], axis=1)
2 df_steps
```

	age	age_cuts	wage
0	18	(17.938, 33.5]	18
1	24	(17.938, 33.5]	24
2	45	(33.5, 49.0]	45
3	43	(33.5, 49.0]	43
4	50	(49.0, 64.5]	50
—	—	—	—
2995	44	(33.5, 49.0]	44
2996	30	(17.938, 33.5]	30
2997	27	(17.938, 33.5]	27
2998	27	(17.938, 33.5]	27
2999	55	(49.0, 64.5]	55

3000 rows × 3 columns

Step Function

Wage Example

```
1 # Create dummy variables for the age groups
2 X = pd.get_dummies(df_steps, drop_first = True).drop(['wage', 'age'],1)
3 X
```

	age_outs_(33.5, 49.0]	age_outs_(49.0, 64.5]	age_outs_(64.5, 80.0]
0	0	0	0
1	0	0	0
2	1	0	0
3	1	0	0
4	0	1	0
—	—	—	—
2995	1	0	0
2996	0	0	0
2997	0	0	0
2998	0	0	0
2999	0	1	0

3000 rows × 3 columns

Step Function

Wage Example

```
1 bin_mapping = np.digitize(ages, bins)
2 bin_mapping

array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3,
       3, 3, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4], dtype=int64)
```

```
1 ages_steps = pd.get_dummies(bin_mapping).drop(1, axis=1)
2 ages_steps
```

	2	3	4
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0
...
57	0	0	1
58	0	0	1
59	0	0	1
60	0	0	1
61	0	0	1


62 rows × 3 columns

Step Function

Wage Example

```
1 lregmodel = LinearRegression(fit_intercept = True)
2 lregmodel.fit(X,y = wage['wage'])
3 y_hat = lregmodel.predict(ages_steps)
4 sns.set(rc = {'figure.figsize':(15,8)})
5 sns.scatterplot(data = wage, x = 'age', y = 'wage')
6 sns.lineplot(x=ages, y=y_hat, color = 'r')
```

<AxesSubplot::xlabel='age', ylabel='wage'>



Step Functions - Summary

- Special case of a basis function where:

$$b_j(X_1) = I(c_j \leq X_1 \leq c_{j+1})$$

- Easy to work with and interpret
- Works well when there are natural cutpoints (or *knots*) to work with
- The discontinuous step function can introduce significant issues

Piecewise Polynomials

Overview



Instead of a single polynomial in X over its whole domain, we can rather use different polynomials in regions defined by knots. For example, a piecewise cubic polynomial with one knot at c :

$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x_i < c \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & \text{if } x_i \geq c \end{cases}$$

- In general, we also add constraints to the polynomials for continuity



Piecewise Polynomial Examples

Discontinuity
at the cut-
point is
undesirable



Adding a
constraint that
curve must be
continuous

Adding
constraints that
first and second
derivatives must
be continuous



Piecewise
polynomial of
order 1

Piecewise Polynomials

Overview

- Piecewise polynomials improve on the "global" polynomial regression because we are able to use lower-order polynomials resulting in less risk of overfitting
 - *Note: the step function is thus a piecewise polynomial of degree 0*

Linear Splines


A linear spline with knots at $\xi_k, k = 1, \dots, K$ is a piecewise linear polynomial continuous at each knot

- We can represent this model with basis functions:

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+1} b_{K+1}(x_i) + \epsilon_i$$


where the b_k functions are basis functions:

$$\begin{aligned} b_1(x_i) &= x_i \\ b_{k+1}(x_i) &= (x_i - \xi_k)_+, \quad k = 1, \dots, K \end{aligned}$$




“Positive part”: returns 0 if value is negative


Linear Splines




Linear Splines



Linear Splines



Piecewise Linear Regression Models



Cubic Splines

A linear cubic with knots at $\xi_k, k = 1, \dots, K$ is a piecewise cubic polynomial with continuous derivatives up to order 2 at each knot

- Again, we can represent this model with truncated basis functions:

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i$$

where:

$$\begin{aligned} b_1(x_i) &= x_i \\ b_2(x_i) &= x_i^2 \\ b_3(x_i) &= x_i^3 \\ b_{k+3}(x_i) &= (x_i - \xi_k)_+^3, \quad k = 1, \dots, K \end{aligned}$$

Cubic Splines

Wage Example

```
1 from statsmodels.gam.api import GLMGam, BSplines
2 bs = BSplines(wage[['age']], df=7, degree=3)
3 reg = GLMGam.from_formula('wage ~ age', wage, smoother=bs).fit()
4 reg.summary()
```

Generalized Linear Model Regression Results


Dep. Variable:	wage	No. Observations:	3000
Model:	GLMGam	Df Residuals:	2993
Model Family:	Gaussian	Df Model:	6.00
Link Function:	Identity	Scale:	1592.5
Method:	PIRLS	Log-Likelihood:	-15313.
Date:	Sun, 31 Jul 2022	Deviance:	4.7662e+06
Time:	15:10:45	Pearson chi2:	4.77e+06
No. iterations:	3		
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
Intercept	36.1720	9.878	3.662	0.000	16.812	55.532
age	1.1190	0.161	6.945	0.000	0.803	1.435
age_s0	21.9493	11.786	1.862	0.063	-1.150	45.049
age_s1	39.2359	5.849	6.708	0.000	27.772	50.700
age_s2	38.6929	5.676	6.817	0.000	27.569	49.817
age_s3	11.4261	6.444	1.773	0.076	-1.204	24.056
age_s4	13.5710	12.406	1.094	0.274	-10.745	37.887
age_s5	-54.6265	10.772	-5.071	0.000	-75.739	-33.514

Cubic Splines

Wage Example



```
1 reg.plot_partial(0, cpr=True, plot_se=False)
```



Regression Splines

- Regression splines are defined to be piecewise polynomials there are continuous at the *knots* (cut-points). That is, the end points where segments meet are constrained to be equal (upper right)
- We can also add constraints that the first and second derivatives at the end points meet are also constrained to be equal, giving a smooth curve

Cubic Splines – Knot Selection



Knot Placement

- One strategy is to decide K , the number of knots, and then place them at appropriate quantiles of the observed X
- A cubic spline with K knots has $K+4$ parameters (or degrees of freedom)

Generalized Additive Models



Generalized Additive Models

Overview

- Extension of multiple linear regression models which replaces each linear component $(\beta_j x_{ij})$ with a smooth nonlinear function $f_j(x_{ij})$.
- Linear regression model then is generalized from this:

$$y_i = \beta_0 + \sum_{j=1}^p \beta_j x_{ij} + \epsilon_i$$

to this:

$$y_i = \beta_0 + \sum_{j=1}^p f_j(x_{ij}) + \epsilon_i$$

Generalized Additive Models

Overview

- Allows combination of regression splines, smoothing splines, and local regression models to deal with multiple predictors in one model
 - Modeler can decide which method to use for every feature

General Additive Models

Extending and Generalizing to Multi-Variable Models

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i.$$

